


1. The simple threshold classifier for sentiment analysis described in the video (*check all that apply*): 1 point



- ☒ Must have pre-defined positive and negative attributes
- ☐ Must either count attributes equally or pre-define weights on attributes
- ☐ Defines a possibly non-linear decision boundary

2. For a linear classifier classifying between “positive” and “negative” sentiment in a review  $x$ ,  $\text{Score}(x) = 0$  implies (*check all that apply*): 1 point

- ☐ The review is very clearly “negative”
- ☒ We are uncertain whether the review is “positive” or “negative”
- ☐ We need to retrain our classifier because an error has occurred

3. For which of the following datasets would a linear classifier perform perfectly? 1 point

☐



4. *True or false:* High classification accuracy always indicates a good classifier.

1 point



True



False

5. *True or false:* For a classifier classifying between 5 classes, there always exists a classifier with accuracy greater than 0.18. 1 point

- ☒ True  
☐ False

6. *True or false:* A false negative is always worse than a false positive. 1 point

- ☐ True  
☒ False

7. Which of the following statements are true? (*Check all that apply*) 1 point

- ☒ Test error tends to decrease with more training data until a point, and then does not change (i.e., curve flattens out)  
☐ Test error always goes to 0 with an unboundedly large training dataset  
☐ Test error is never a function of the amount of training data